



Institute of Distance and Open Learning

Vidya Nagari, Kalina, Santacruz East – 400098.

A Practical Journal Submitted in fulfillment of the degree of

MASTER OF SCIENCE IN COMPUTER SCIENCE

YEAR 2024-25

Part - II

Semester - IV

Subject Code- 91955

Subject Name- Business Intelligence and Big Data Analytics III

BY

RAINA RIZWANULLAH KHAN

Seat Number: 2700155

Application ID: 42662

Institute of Distance and Open Learning (IDOL)

University of Mumbai



Certificate

This is to certify that **Miss Raina Rizwanullah Khan** student of Masters of Computer Science, Part 2, Semester 4 has completed the specified term work in the subject of **Business Intelligence and Big Data Analytics III** in satisfactorily manner within this institute as laid down by University of Mumbai during the academic year 2024 to 2025.

M.Sc. CS Coordinator

Examiner

Date:30-06-2025

Guide

INDEX

No.	Practical Name	Page No.
1	Pre-process the given data set and hence apply clustering techniques like KMeans, K-Medoids. Interpret the result.	04-08
2	Pre-process the given data set and hence apply partition clustering algorithms. Interpret the result	9-10
3	Pre-process the given data set and hence apply hierarchical algorithms and density-based clustering techniques. Interpret the result.	11-15
4	Pre-process the given data set and hence classify the resultant data set using tree classification techniques. Interpret the result.	16-19
5	Pre-process the given data set and hence classify the resultant data set using Statistical based classifiers. Interpret the result.	20-21
6	Pre-process the given data set and hence classify the resultant data set using support vector machine. Interpret the result.	22-24
7	Write a program to explain different functions of Principal Components.	25-27

Practical No 1

Aim: Pre-process the given data set and hence apply clustering techniques like KMeans, K-Medoids. Interpret the result.

Code:

Install python on the system then run this code in cmd `python -m pip install pandas numpy matplotlib scikit-learn scikit-learn-extra` to install all the required packages.

```
C:\Users\Raina Khan>python -m pip uninstall numpy scikit-learn-extra -y
Found existing installation: numpy 2.2.6
Uninstalling numpy-2.2.6:
  Successfully uninstalled numpy-2.2.6
```

```
C:\Users\Raina Khan>python -m pip install numpy==1.23.5
Collecting numpy==1.23.5
  Downloading numpy-1.23.5-cp310-cp310-win_amd64.whl (14.6 MB)
    | 14.6 MB 2.2 MB/s
Installing collected packages: numpy
Successfully installed numpy-1.23.5
WARNING: You are using pip version 21.2.3; however, version 25.1.1 is available.
You should consider upgrading via the 'C:\Users\Raina Khan\AppData\Local\Programs\Python\Python310\python.exe -m pip install --upgrade pip' command.
```

```
C:\Users\Raina Khan>python -m pip install scikit-learn-extra --no-binary :all:
```

```
C:\Users\Raina Khan>python -m pip install --upgrade pip
Requirement already satisfied: pip in c:\users\raina khan\appdata\local\programs\python\python310\lib\site-packages (21.2.3)
Collecting pip
  Downloading pip-25.1.1-py3-none-any.whl (1.8 MB)
    | 1.8 MB 3.2 MB/s
Installing collected packages: pip
  Attempting uninstall: pip
    Found existing installation: pip 21.2.3
    Uninstalling pip-21.2.3:
      Successfully uninstalled pip-21.2.3
  Successfully installed pip-25.1.1
```

Install required packages first via Command Prompt (if not installed):

`python -m pip install pandas numpy matplotlib scikit-learn scikit-learn-extra`

--- Start of Code ---

`import pandas as pd`

`import numpy as np`

`from sklearn.preprocessing import StandardScaler`

`from sklearn.cluster import KMeans`

`from sklearn_extra.cluster import KMedoids`

`from sklearn.decomposition import PCA`

```

import matplotlib.pyplot as plt
from sklearn.datasets import make_blobs

# 1. Generate synthetic dataset (150 samples, 4 features, 3 clusters)
X, _ = make_blobs(n_samples=150, centers=3, n_features=4,
random_state=42)
df = pd.DataFrame(X, columns=['Feature1', 'Feature2', 'Feature3', 'Feature4'])

# 2. Pre-processing: Standardize the data
scaler = StandardScaler()
df_scaled = scaler.fit_transform(df)

# 3. Apply KMeans clustering
kmeans = KMeans(n_clusters=3, random_state=42)
kmeans.fit(df_scaled)
df['KMeans_Cluster'] = kmeans.labels_

# Output KMeans results
print("KMeans Cluster Centers:\n", kmeans.cluster_centers_)

# 4. Apply KMedoids clustering
kmedoids = KMedoids(n_clusters=3, random_state=42)
kmedoids.fit(df_scaled)
df['KMedoids_Cluster'] = kmedoids.labels_

# Output KMedoids results
print("KMedoids Cluster Medoids (indices):", kmedoids.medoid_indices_)

```

5. Reduce dimensions for visualization

```
pca = PCA(n_components=2)
df_pca = pca.fit_transform(df_scaled)
```

6. Visualize KMeans clusters

```
plt.figure(figsize=(8, 6))
plt.scatter(df_pca[:, 0], df_pca[:, 1], c=df['KMeans_Cluster'], cmap='viridis',
alpha=0.7)
plt.title('KMeans Clustering (PCA Projection)')
plt.xlabel('PCA Component 1')
plt.ylabel('PCA Component 2')
plt.colorbar(label='Cluster')
plt.grid(True)
plt.tight_layout()
plt.show()
```

7. Visualize KMedoids clusters

```
plt.figure(figsize=(8, 6))
plt.scatter(df_pca[:, 0], df_pca[:, 1], c=df['KMedoids_Cluster'], cmap='plasma',
alpha=0.7)
plt.title('KMedoids Clustering (PCA Projection)')
plt.xlabel('PCA Component 1')
plt.ylabel('PCA Component 2')
plt.colorbar(label='Cluster')
plt.grid(True)
plt.tight_layout()
```

```
plt.show()
```

```
# Optional: Save result
```

```
# df.to_csv("clustered_result.csv", index=False)
```

```
# --- End of Code ---
```

Output:





Figure 1



```
= RESTART: C:/Users/Raina Khan/AppData/Local/Programs/Python/Python310/BI Pract
1.py
KMeans Cluster Centers:
[[-1.16516133 -1.33142953 -0.62509016  0.34036025]
 [-0.04419554  1.03916665  1.3942501  -1.29887211]
 [ 1.20935687  0.29226288 -0.76915994  0.95851186]]
KMedoids Cluster Medoids (indices): [133 144  48]
```


Practical No 2

Aim: Pre-process the given data set and hence apply partition clustering algorithms. Interpret the result.

Code:

```
import pandas as pd
import numpy as np
from sklearn.datasets import make_blobs
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans

# Step 1: Generate synthetic dataset
X, _ = make_blobs(n_samples=100, centers=3, n_features=3,
random_state=42)
df = pd.DataFrame(X, columns=['Feature1', 'Feature2', 'Feature3'])

# Step 2: Preprocess (scale) the data
scaler = StandardScaler()
scaled_features = scaler.fit_transform(df)
df_scaled = pd.DataFrame(scaled_features, columns=df.columns)

# Step 3: Apply K-Means clustering
num_clusters = 3
kmeans = KMeans(n_clusters=num_clusters, random_state=42, n_init=10)
kmeans.fit(df_scaled)

# Step 4: Add cluster labels
df['Cluster'] = kmeans.labels_
```

Step 5: Print number of data points in each cluster

```
print("\nNumber of data points in each cluster:")
```

```
print(df['Cluster'].value_counts())
```

Step 6: Print cluster centroids

```
print("\nCluster centroids:")
```

```
centroids = pd.DataFrame(kmeans.cluster_centers_, columns=df.columns[:-1])
```

```
print(centroids)
```

Step 7: Print mean values of features across clusters

```
print("\nMean values of features across clusters:")
```

```
cluster_means = df.groupby('Cluster').mean(numeric_only=True)
```

```
print(cluster_means)
```

Output:

```
= RESTART: C:/Users/Raina Khan/AppData/Local/Programs/Python/Python310/BI Pract2.py
```

```
[[95mNumber of data points in each cluster:[0m
```

```
Cluster
```

```
1      34
```

```
2      33
```

```
0      33
```

```
Name: count, dtype: int64
```

```
[[95mCluster centroids:[0m
```

```
   Feature1  Feature2  Feature3
```

```
0  1.141480 -1.407411 -1.358115
```

```
1  0.110490  0.789528  0.929131
```

```
2 -1.255318  0.593958  0.400828
```

```
[[95mMean values of features across clusters:[0m
```

```
   Feature1  Feature2  Feature3
```

```
Cluster
```

```
0    1.951547 -6.481144 -6.634443
```

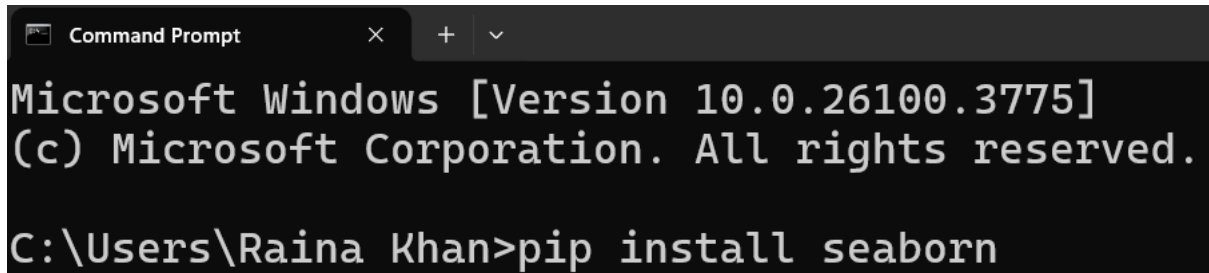
```
1   -2.645974  8.887457  4.697897
```

```
2   -8.736560  7.519361  2.080378
```

Practical No 3

Aim: Pre-process the given data set and hence apply hierarchical algorithms and density-based clustering techniques. Interpret the result.

Code:

A screenshot of a Windows Command Prompt window. The title bar says 'Command Prompt'. The window content shows the Microsoft Windows version (10.0.26100.3775) and copyright information. The command prompt shows the user 'Raina Khan' at the 'C:\Users\' directory, and the command 'pip install seaborn' has been entered.

```
Microsoft Windows [Version 10.0.26100.3775]
(c) Microsoft Corporation. All rights reserved.

C:\Users\Raina Khan>pip install seaborn
```

```
import pandas as pd
import numpy as np
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import DBSCAN
from scipy.cluster.hierarchy import linkage, dendrogram, fcluster
import matplotlib.pyplot as plt
import seaborn as sns

# Optional: Set style for plots
sns.set(style="whitegrid")

# Step 1: Load dataset
# If you don't have a file, you can generate synthetic data using make_blobs
# (see below)
# df = pd.read_csv("your_dataset.csv")
from sklearn.datasets import make_blobs
X, _ = make_blobs(n_samples=100, centers=3, n_features=3,
random_state=42)
```

```
df = pd.DataFrame(X, columns=["Feature1", "Feature2", "Feature3"])
```

```
# Step 2: Handle missing values (drop rows with NaNs)
```

```
df.dropna(inplace=True)
```

```
# Step 3: Scale the data
```

```
scaler = StandardScaler()
```

```
scaled_data = scaler.fit_transform(df)
```

```
# -----
```

```
# Step 4: Hierarchical Clustering
```

```
# -----
```

```
linked = linkage(scaled_data, method='ward')
```

```
# Plot the dendrogram
```

```
plt.figure(figsize=(10, 6))
```

```
dendrogram(linked, truncate_mode='lastp', p=30, leaf_rotation=45.,
```

```
leaf_font_size=10., show_contracted=True)
```

```
plt.title('Hierarchical Clustering Dendrogram')
```

```
plt.xlabel('Sample Index')
```

```
plt.ylabel('Distance')
```

```
plt.show()
```

```
# Assign cluster labels (e.g., k=3)
```

```
hc_labels = fcluster(linked, t=3, criterion='maxclust')
```

```
df['HC_Cluster'] = hc_labels
```

```

# -----
# Step 5: DBSCAN Clustering
# -----

db = DBSCAN(eps=0.5, min_samples=5)
db_labels = db.fit_predict(scaled_data)
df['DBSCAN_Cluster'] = db_labels


# -----
# Step 6: Visualize Clusters
# -----

# Visualize Hierarchical Clusters
plt.figure(figsize=(8, 5))
sns.scatterplot(data=df, x="Feature1", y="Feature2", hue="HC_Cluster",
palette="Set2")
plt.title("Hierarchical Clustering Result")
plt.show()


# Visualize DBSCAN Clusters
plt.figure(figsize=(8, 5))
sns.scatterplot(data=df, x="Feature1", y="Feature2", hue="DBSCAN_Cluster",
palette="Set1")
plt.title("DBSCAN Clustering Result")
plt.show()


# -----
# Step 7: Interpretation Guide
# -----

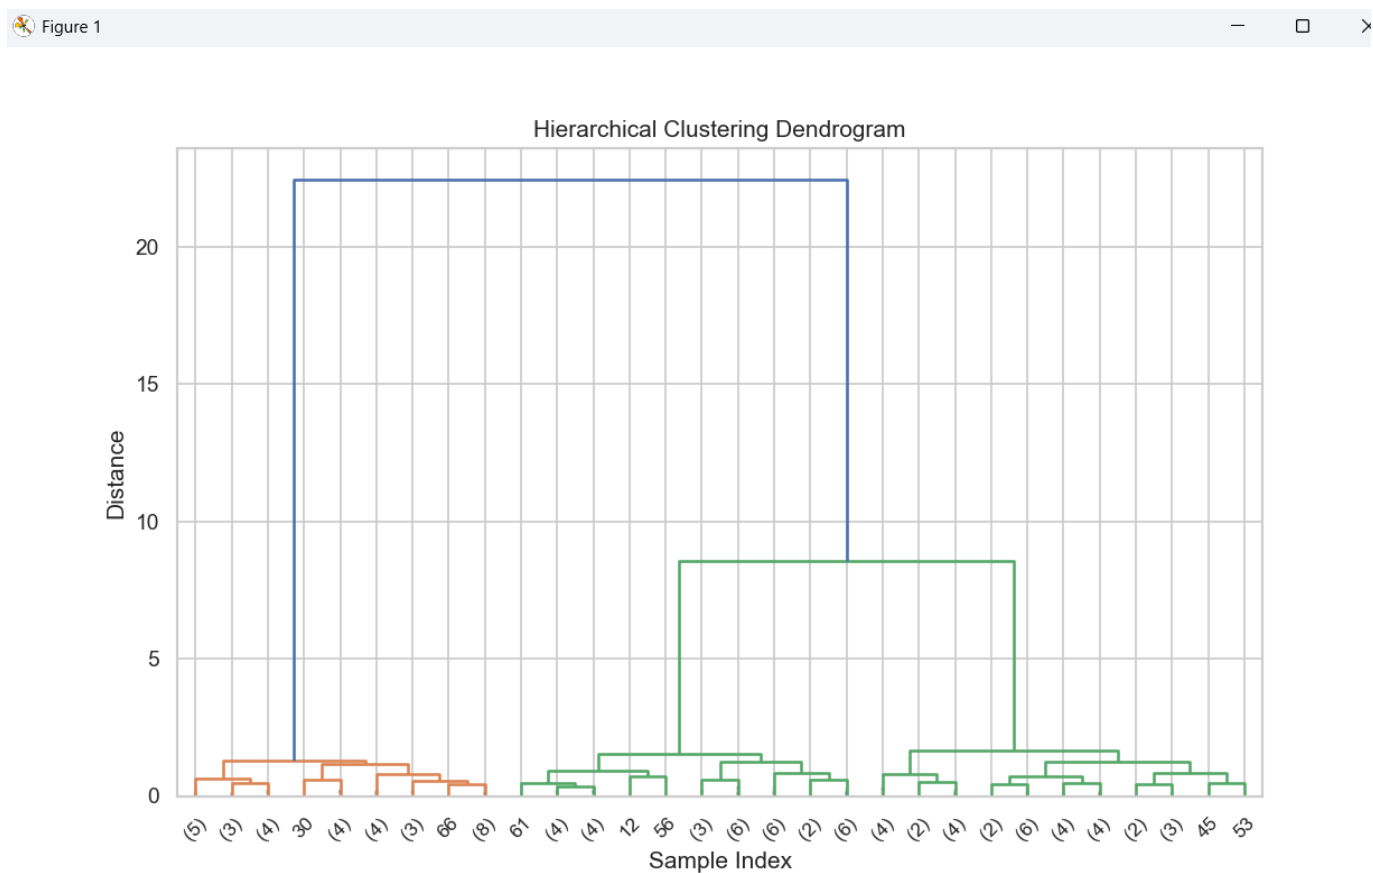
```

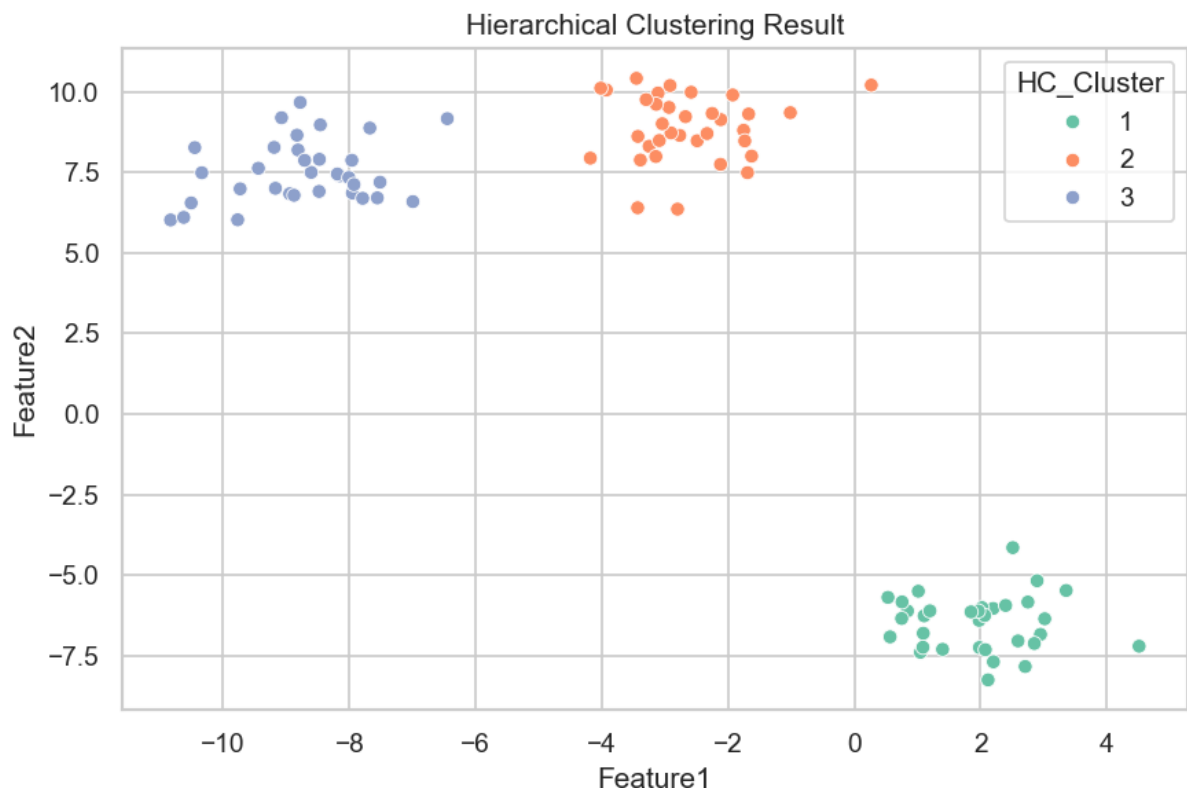
```

print("\nInterpretation Guide:")
print("- Hierarchical clustering builds a tree (dendrogram). Cut the tree to
decide number of clusters.")
print("- DBSCAN identifies clusters based on density. Noise points (outliers) are
labeled as -1.")
print("- You can adjust `eps` and `min_samples` to tune DBSCAN for your
dataset.")

```

Output:





Practical No 4

Aim: Pre-process the given data set and hence classify the resultant data set using tree classification techniques. Interpret the result.

Code:

```
# Import necessary libraries
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, accuracy_score
import matplotlib.pyplot as plt

# Step 1: Generate a synthetic dataset (4 features, 2 classes)
from sklearn.datasets import make_classification

X_sample, y_sample = make_classification(
    n_samples=220,
    n_features=4,
    n_informative=3,
    n_redundant=0,
    n_repeated=0,
    n_classes=2,
    random_state=42
)

# Step 2: Convert to DataFrame
```



```

df = pd.DataFrame(X_sample, columns=["Feature1", "Feature2", "Feature3",
"Feature4"])
df["target"] = y_sample

# Step 3: Preprocessing (optional: drop NA if real dataset is used)
df.dropna(inplace=True)
le = LabelEncoder()
df["target"] = le.fit_transform(df["target"])

# Step 4: Split dataset into features and labels
X = df.drop("target", axis=1)
y = df["target"]
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)

# -----
# Decision Tree Classifier
# -----
dt_classifier = DecisionTreeClassifier(random_state=42)
dt_classifier.fit(X_train, y_train)

# Evaluate
dt_predictions = dt_classifier.predict(X_test)
dt_accuracy = accuracy_score(y_test, dt_predictions)
print("Decision Tree Accuracy:", round(dt_accuracy, 2))
print(classification_report(y_test, dt_predictions))

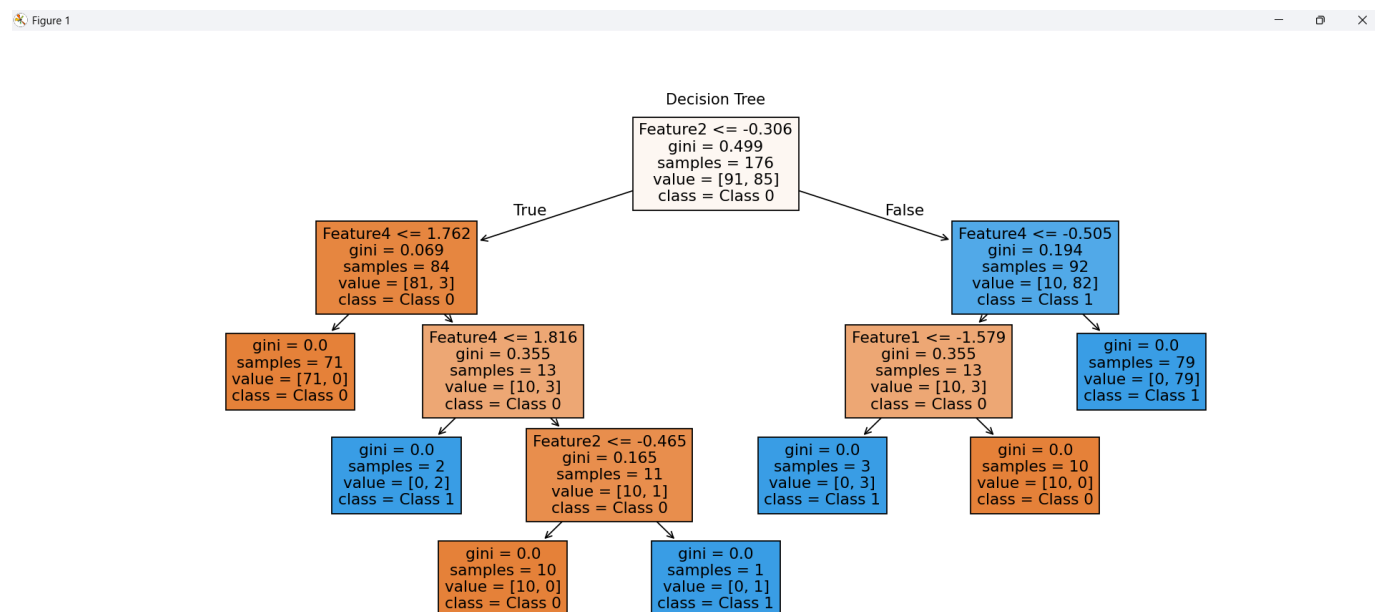
```

```
# Plot Decision Tree
plt.figure(figsize=(12, 8))
plot_tree(dt_classifier, filled=True, feature_names=X.columns,
class_names=["Class 0", "Class 1"])
plt.title("Decision Tree")
plt.show()

# -----
# Random Forest Classifier
# -----
rf_classifier = RandomForestClassifier(random_state=42)
rf_classifier.fit(X_train, y_train)

# Evaluate
rf_predictions = rf_classifier.predict(X_test)
rf_accuracy = accuracy_score(y_test, rf_predictions)
print("Random Forest Accuracy:", round(rf_accuracy, 2))
print(classification_report(y_test, rf_predictions))
```

Output:



= RESTART: C:/Users/Raina Khan/AppData/Local/Programs/Python/Python310/BI Pract 4.py

Decision Tree Accuracy: 0.93

	precision	recall	f1-score	support
0	0.90	0.95	0.93	20
1	0.96	0.92	0.94	24
accuracy			0.93	44
macro avg	0.93	0.93	0.93	44
weighted avg	0.93	0.93	0.93	44

Random Forest Accuracy: 0.93

	precision	recall	f1-score	support
0	0.95	0.90	0.92	20
1	0.92	0.96	0.94	24
accuracy			0.93	44
macro avg	0.93	0.93	0.93	44
weighted avg	0.93	0.93	0.93	44

Practical No 5

Aim: Pre-process the given data set and hence classify the resultant data set using Statistical based classifiers. Interpret the result.

Code:

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import classification_report, accuracy_score
from sklearn.datasets import make_classification

# Generate synthetic dataset
X, y = make_classification(
    n_samples=250,
    n_features=5,
    n_informative=3,
    n_redundant=0,
    n_classes=2,
    random_state=42
)

# Convert to DataFrame for consistency
df = pd.DataFrame(X, columns=[f'Feature{i}' for i in range(1, 6)])
df['target'] = y

# Step 1: Separate features and target
X = df.drop('target', axis=1)
```

```

y = df['target']

# Step 2: Normalize features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Step 3: Train-test split
X_train, X_test, y_train, y_test = train_test_split(
    X_scaled, y, test_size=0.2, random_state=42
)

# Step 4: Train Naive Bayes classifier
classifier = GaussianNB()
classifier.fit(X_train, y_train)

# Step 5: Evaluate
y_pred = classifier.predict(X_test)
print("Classification Report:")
print(classification_report(y_test, y_pred))
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy}")

```

Output:

```

= RESTART: C:/Users/Raina Khan/AppData/Local/Programs/Python/Python310/BI Pract
5.py
Classification Report:

```

	precision	recall	f1-score	support
0	0.92	0.85	0.88	26
1	0.85	0.92	0.88	24
accuracy			0.88	50
macro avg	0.88	0.88	0.88	50
weighted avg	0.88	0.88	0.88	50

```

Accuracy: 0.88

```

Practical No 6

Aim: Pre-process the given data set and hence classify the resultant data set using support vector machine. Interpret the result

Code:

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.svm import SVC
from sklearn.metrics import classification_report, confusion_matrix
import matplotlib.pyplot as plt

# Step 1: Generate synthetic dataset (replace this with pd.read_csv if you have
a real file)
from sklearn.datasets import make_classification
X, y = make_classification(
    n_samples=250,
    n_features=2, # Set to 2 for plotting decision boundary
    n_informative=2,
    n_redundant=0,
    n_classes=2,
    random_state=42
)

# Step 2: Scale the features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

```

# Step 3: Split into training and testing
X_train, X_test, y_train, y_test = train_test_split(
    X_scaled, y, test_size=0.2, random_state=42
)

# Step 4: Train SVM (linear kernel)
svm_model = SVC(kernel='linear')
svm_model.fit(X_train, y_train)

# Step 5: Predict & Evaluate
y_pred = svm_model.predict(X_test)
print("Classification Report:")
print(classification_report(y_test, y_pred))
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))

# Step 6: Plot if 2D
if X_train.shape[1] == 2:
    plt.figure(figsize=(8, 6))
    plt.scatter(X_train[:, 0], X_train[:, 1], c=y_train, cmap='viridis',
edgecolors='k')
    plt.xlabel('Feature 1')
    plt.ylabel('Feature 2')

# Plot decision boundary
ax = plt.gca()


```

```

xlim = ax.get_xlim()
ylim = ax.get_ylim()
xx, yy = np.meshgrid(np.linspace(xlim[0], xlim[1], 50),
                      np.linspace(ylim[0], ylim[1], 50))
Z = svm_model.decision_function(np.c_[xx.ravel(), yy.ravel()])
Z = Z.reshape(xx.shape)
ax.contour(xx, yy, Z, colors='k', levels=[-1, 0, 1],
           alpha=0.5, linestyles=['--', '-', '--'])
plt.title("SVM Decision Boundary")
plt.show()

```

Output:



= RESTART: C:/Users/Raina Khan/AppData/Local/Programs/Python/Python310/BI Pract 6.py

Classification Report:

	precision	recall	f1-score	support
0	0.74	0.74	0.74	19
1	0.84	0.84	0.84	31
accuracy			0.80	50
macro avg	0.79	0.79	0.79	50
weighted avg	0.80	0.80	0.80	50

Confusion Matrix:

```
[[14  5]
 [ 5 26]]
```


Practical No 7

Aim: Write a program to explain different functions of Principal Components.

Code:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler

# Step 1: Create a sample dataset with more meaningful variance
X = np.array([
    [2.5, 2.4],
    [0.5, 0.7],
    [2.2, 2.9],
    [1.9, 2.2],
    [3.1, 3.0],
    [2.3, 2.7],
    [2, 1.6],
    [1, 1.1],
    [1.5, 1.6],
    [1.1, 0.9]
])

print("Original Data:\n", X)

# Step 2: Standardize the data (very important for PCA)
scaler = StandardScaler()
```

```
X_scaled = scaler.fit_transform(X)
```

```
# Step 3: Apply PCA
```

```
pca = PCA(n_components=2) # Keep both components for explanation
```

```
X_pca = pca.fit_transform(X_scaled)
```

```
# Step 4: Display results
```

```
print("\nTransformed Data (PCA Result):\n", X_pca)
```

```
print("\nExplained Variance Ratio:", pca.explained_variance_ratio_)
```

```
print("Singular Values:", pca.singular_values_)
```

```
print("Components (eigenvectors):\n", pca.components_)
```

```
# Step 5: Visualize the PCA result
```

```
plt.figure(figsize=(8, 6))
```

```
plt.scatter(X_pca[:, 0], X_pca[:, 1], color='green', edgecolor='k')
```

```
plt.title('PCA Transformed Data')
```

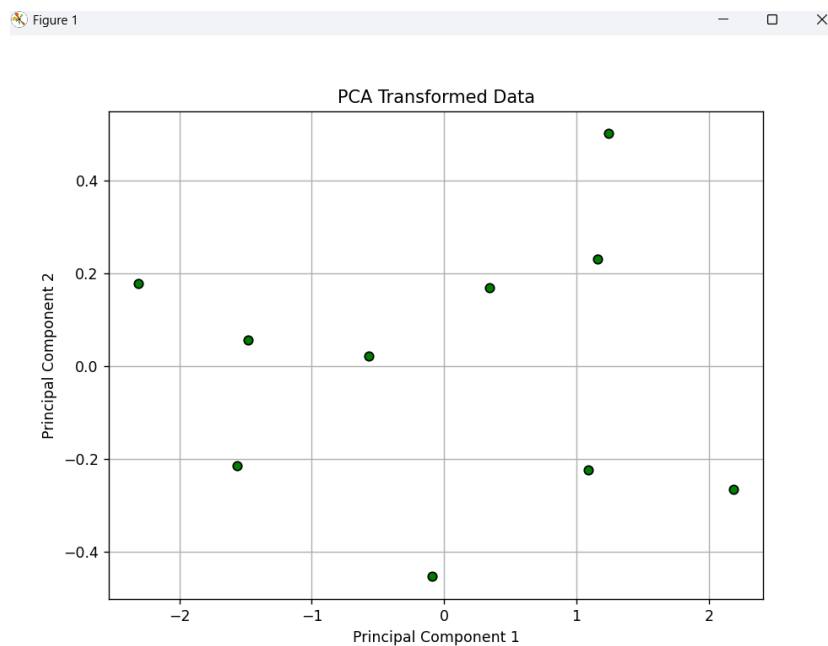
```
plt.xlabel('Principal Component 1')
```

```
plt.ylabel('Principal Component 2')
```

```
plt.grid(True)
```

```
plt.show()
```

Output:



```
= RESTART: C:/Users/Raina Khan/AppData/Local/Programs/Python/Python310/BI Pract 7.py
```

```
Original Data:
```

```
[[2.5 2.4]
 [0.5 0.7]
 [2.2 2.9]
 [1.9 2.2]
 [3.1 3. ]
 [2.3 2.7]
 [2.  1.6]
 [1.  1.1]
 [1.5 1.6]
 [1.1 0.9]]
```

```
Transformed Data (PCA Result):
```

```
[[ 1.08643242 -0.22352364]
 [-2.3089372  0.17808082]
 [ 1.24191895  0.501509 ]
 [ 0.34078247  0.16991864]
 [ 2.18429003 -0.26475825]
 [ 1.16073946  0.23048082]
 [-0.09260467 -0.45331721]
 [-1.48210777  0.05566672]
 [-0.56722643  0.02130455]
 [-1.56328726 -0.21536146]]
```

```
Explained Variance Ratio: [0.96296464 0.03703536]
```

```
Singular Values: [4.38854107 0.86064352]
```

```
Components (eigenvectors):
```

```
[[ 0.70710678  0.70710678]
 [-0.70710678  0.70710678]]
```